

Video link: <https://drive.google.com/file/d/1ROnWeKZzW1d2vbceApyy9kmipLMf1GNn/view?usp=sharing>

Github link: <https://github.com/naraanjali/Assignment-4>

The screenshot displays a Jupyter Notebook titled 'Anjali_Assig4.ipynb' in a dark-themed IDE. The notebook is open to a cell containing Python code for data analysis. The code imports pandas and matplotlib, reads a CSV file, and performs various statistical operations. The output of the code is visible in the console below the code cell.

```
#Import Libraries
import pandas as pd
import matplotlib.pyplot as plt
df = pd.read_csv('data.csv') #Read CSV File
print("Basic Statistical Description:")
print(df.describe()) #Basic Statistical Description
print("\nCheck for Null Values:") #Check for Null Values
print(df.isnull().sum())
df.fillna(df.mean(), inplace=True) #Fill Null Values
agg_columns = ['Calories', 'Duration'] #Aggregate Data
agg_result = df[agg_columns].agg(['min', 'max', 'count', 'mean'])
print("\nAggregated Data:")
print(agg_result)
filtered_df1 = df[(df['Calories'] >= 500) & (df['Calories'] <= 1000)] #Filter Data
filtered_df2 = df[(df['Calories'] > 500) & (df['Pulse'] < 100)]
df_modified = df.drop(columns=['Maxpulse']) #Modify DataFrame
df.drop(columns=['Maxpulse'], inplace=True)
df['Calories'] = df['Calories'].astype(int) #Data Type Conversion
df.plot.scatter(x='Duration', y='Calories', title='Scatter Plot: Duration vs Calories') #Scatter Plot
plt.show()
```


The console output shows the following results:

```
Basic Statistical Description:
count    169.000000    169.000000    169.000000    164.000000
mean      63.846154    107.461538    134.047337    375.790244
std       42.299949    14.510259    16.450434    266.379919
min       15.000000     80.000000    100.000000     50.300000
25%       45.000000    100.000000    124.000000    250.925000
50%       60.000000    105.000000    131.000000    318.600000
75%       60.000000    111.000000    141.000000    387.600000
max       300.000000    159.000000    184.000000    1860.400000

Check for Null Values:
Duration    0
Pulse       0
Maxpulse    0
Calories    5
dtype: int64

Aggregated Data:
      Calories  Duration
min    50.300000    15.000000
max   1860.400000   300.000000
count   169.000000   169.000000
mean    375.790244    63.846154
```

Below the console output, a scatter plot titled 'Scatter Plot: Duration vs Calories' is displayed. The plot shows a positive correlation between Duration (x-axis) and Calories (y-axis). The x-axis ranges from 15 to 300, and the y-axis ranges from 50 to 1860. The plot area is currently blank, suggesting the plot may not have rendered correctly or is partially obscured.



File Edit Selection View Go Run ... Search

Anjali_Assig4.ipynb

C: > ALL > photos > ALL_FILES > CourseWork > Neural_Networks > Anjali_Assig4.ipynb > #Import Libraries

+ Code + Markdown ▶ Run All ⌂ Restart ☰ Clear All Outputs | Variables Outline ...

base (Python 3.11.5)

```
#Import Libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt

# Load the Dataset"
data = pd.read_csv("Salary_Data.csv")
print("Frist few rows of the data frame")
print(data.head())

# b) Split the data into train_test partitions
X = data[['YearsExperience']] # Assuming the independent variable is in the 'YearsExperience' column
y = data['Salary']           # Assuming the dependent variable is in the 'Salary' column

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42)

#Train and predict the model
model = LinearRegression()
model.fit(X_train, y_train)
y_train_pred = model.predict(X_train)
y_test_pred = model.predict(X_test)

#Calculate the mean_squared error
mse_train = mean_squared_error(y_train, y_train_pred)
mse_test = mean_squared_error(y_test, y_test_pred)

print(f"Mean Squared Error (Train): {mse_train}")
```

Ln 25, Col 54 3.12.1 64-bit Cell 3 of 3 Go Live 18:53 02-02-2024

File Edit Selection View Go Run ... Search

Anjali_Assig4.ipynb

C: > ALL > photos > ALL_FILES > CourseWork > Neural_Networks > Anjali_Assig4.ipynb > #Import Libraries

+ Code + Markdown ▶ Run All ⌂ Restart ☰ Clear All Outputs 📄 Variables 📄 Outline ...

base (Python 3.11.5)

```
y_train_pred = model.predict(X_train)
y_test_pred = model.predict(X_test)

#Calculate the mean_squared error
mse_train = mean_squared_error(y_train, y_train_pred)
mse_test = mean_squared_error(y_test, y_test_pred)

print(f"Mean Squared Error (Train): {mse_train}")
print(f"Mean Squared Error (Test): {mse_test}")

#Visualize both train and test data using scatter plot
plt.scatter(X_train, y_train, color='blue', label='Train Data')
plt.scatter(X_test, y_test, color='red', label='Test Data')
plt.plot(X_train, y_train_pred, color='green', linewidth=2, label='Regression Line')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.title('Salary Prediction Model')
plt.legend()
plt.show()
```

[2] ✓ 1.0s Python

... First few rows of the data frame

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0

Mean Squared Error (Train): 29793161.082422983
Mean Squared Error (Test): 35301898.887134895

Ln 25, Col 54 3.12.1 64-bit Cell 3 of 3 Go Live 18:53 02-02-2024

File Edit Selection View Go Run ... Search

Anjali_Assig4.ipynb

C: > ALL > photos > ALL_FILES > CourseWork > Neural_Networks > Anjali_Assig4.ipynb > #Import Libraries

+ Code + Markdown ▶ Run All ⌂ Restart ☰ Clear All Outputs 📄 Variables 📄 Outline ...

base (Python 3.11.5)

```
3 2.0 43525.0
4 2.2 39891.0

Mean Squared Error (Train): 29793161.082422983
Mean Squared Error (Test): 35301898.887134895
```

...

Salary Prediction Model

Legend: Train Data (blue dots), Test Data (red dots), Regression Line (green line).

Y-axis: Salary (40000 to 120000)

X-axis: Years of Experience (0 to 10)

Ln 25, Col 54 3.12.1 64-bit Cell 3 of 3 Go Live 18:53 02-02-2024